

TRƯỜNG ĐẠI HỌC NAM CẦN THƠ
KHOA CÔNG NGHỆ THÔNG TIN



NGUYỄN THÀNH NGUYỄN

**PHÂN LOẠI BỆNH ĐÁI THÁO ĐƯỜNG TỪ
ẢNH ĐÁY MẮT BẰNG MÔ HÌNH HỌC SÂU**

BÁO CÁO

Ngành: Công nghệ thông tin

Mã số ngành: 7480201

Tháng 6 - năm 2026

TRƯỜNG ĐẠI HỌC NAM CẦN THƠ
KHOA CÔNG NGHỆ THÔNG TIN

NGUYỄN THÀNH NGUYỄN
MSSV: 225255

PHÂN LOẠI BỆNH ĐÁI THÁO ĐƯỜNG TỪ
ẢNH ĐÁY MẮT BẰNG MÔ HÌNH HỌC SÂU

BÁO CÁO

Ngành: Công nghệ thông tin

Mã số ngành: 7480201

GIẢNG VIÊN HƯỚNG DẪN
TRẦN VĂN THIỆN

Tháng 6 - năm 2026

LỜI CẢM TẠ

Chúng em xin gửi lời cảm ơn sâu sắc đến Quý Thầy thuộc Khoa Công nghệ thông tin, Trường Đại học Nam Cần Thơ, vì đã tận tình giảng dạy và truyền đạt những kiến thức chuyên môn quý báu, tạo nền tảng vững chắc cho việc hoàn thành báo cáo thực tập này.

Đồng thời, chúng em đặc biệt tri ân Thầy Trần Văn Thiện vì sự hướng dẫn tận tâm và những hỗ trợ quý giá trong suốt quá trình thực tập, giúp chúng em hoàn thiện bài báo cáo một cách hiệu quả và tốt nhất.

Những kiến thức tiếp thu được trong quá trình học tập kết hợp với kinh nghiệm thực tiễn từ thời gian thực tập không chỉ là cơ sở quan trọng để hoàn thành báo cáo tốt nghiệp, mà còn là hành trang ý nghĩa, giúp tác giả tự tin chinh phục những bước tiến trong sự nghiệp và cuộc sống.

Xin chân thành cảm ơn!

Cần Thơ, ngày ... tháng ... năm 2026

Người thực hiện

NGUYỄN THÀNH NGUYÊN

LỜI CAM KẾT

Tôi xin cam kết rằng báo cáo thực tập tốt nghiệp này được thực hiện dựa trên kết quả nghiên cứu của chính tôi. Toàn bộ nội dung và kết quả nghiên cứu trong báo cáo là hoàn toàn trung thực, chưa từng được sử dụng cho bất kỳ hoạt động nghiên cứu hay công trình nào khác ở cùng cấp.

Tôi chịu hoàn toàn trách nhiệm về tính trung thực và nguyên bản của báo cáo này.

Cần Thơ, ngày ... tháng ... năm 2026

Người thực hiện

NGUYỄN THÀNH NGUYỄN

NHẬN XÉT CỦA GIẢNG VIÊN HƯỚNG DẪN

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

Cần Thơ, ngày tháng năm 2026

Giảng viên hướng dẫn

Trần Văn Thiện

MỤC LỤC

CHƯƠNG 1: INTRODUCTION	1
CHƯƠNG 2: RELATED WORK.....	3
CHƯƠNG 3: DATASET	6
CHƯƠNG 4: THEORETICAL BACKGROUND	9
CHƯƠNG 5: PROPOSITION	11
CHƯƠNG 6: EXPERIMENTATION	13
CHƯƠNG 7: CONCLUSION	15
CHƯƠNG 8: DISCUSSION.....	16
REFERENCES.....	18

DANH MỤC BIỂU BẢNG

Bảng 2. 1. Comparative Analysis of Pioneering Methodologies in Automated Ophthalmic Diagnostics.....	4
Bảng 3. 1. Statistical Distribution and Clinical Characterization of the Target Dataset.	7

DANH MỤC HÌNH

DANH MỤC TỪ VIẾT TẮT

DR	:	Diabetic Retinopathy
NPDR	:	Non-Proliferative Diabetic Retinopathy
PDR	:	Proliferative Diabetic Retinopathy
APTOS	:	Asia Pacific Tele-Ophthalmology Society
CLAHE	:	Contrast Limited Adaptive Histogram Equalization
CDF	:	Cumulative Distribution Function
GAN	:	Generative Adversarial Network
ESRGAN	:	Enhanced Super-Resolution Generative Adversarial Networks
VGG	:	Visual Geometry Group
CNN	:	Convolutional Neural Network
ResNet50	:	Residual Network 50-layers
DenseNet121	:	Densely Connected Convolutional Network 121-layers
ReLU	:	Rectified Linear Unit
CUDA	:	Compute Unified Device Architecture
VRAM	:	Video Random Access Memory
XAI	:	Explainable Artificial Intelligence
Grad-CAM	:	Gradient-weighted Class Activation Mapping
MAE	:	Mean Absolute Error
QWK	:	Quadratic Weighted Kappa
TP	:	True Positive

Explainable Diabetic Retinopathy Classification from Fundus Images Using Deep Learning and Grad-CAM

NGUYEN Thanh-Nguyen¹[0000-1111-2222-3333] and TRAN Van-Thien²[1111-2222-3333-4444]

¹ Faculty of Information Technology, Nam Can Tho University, 168 Nguyen Van Cu Street, An Binh Ward, Can Tho City, Vietnam

{nguyen225255@student.nctu.edu.vn, nhakhoi@nctu.edu.vn}

Abstract

Deep learning models have demonstrated exceptional capabilities in diagnosing diabetic retinopathy (DR) from color fundus photographs. However, clinical adoption remains constrained by image quality degradation in real-world settings and the inherently opaque nature of deep neural networks. This study proposes an integrated clinical decision support system that addresses these limitations through a sequential pipeline of image enhancement, ensemble classification, and visual explanation. First, Contrast Limited Adaptive Histogram Equalization (CLAHE) and Enhanced Super-Resolution Generative Adversarial Networks (ESRGAN) are employed to eliminate capture noise and restore micro-lesion morphological details. Second, a weighted ensemble architecture combining ResNet50 and DenseNet121 is optimized using Focal Loss to mitigate class imbalance across the five severity levels of DR. Finally, Gradient-weighted Class Activation Mapping (Grad-CAM) is integrated to generate visual heatmaps of pathological regions, bridging the gap between mathematical optimization and clinical interpretability. Experimental evaluation on the benchmarking independent test set yields an overall test accuracy of 75.48%, a Quadratic Weighted Kappa of 0.84301, and a Mean Absolute Error of 0.33787, confirming the framework's viability for large-scale automated screening.

Keywords: Diabetic Retinopathy, Deep Learning, Ensemble Learning, Explainable AI, Super-Resolution, Clinical Decision Support System.

CHƯƠNG 1: INTRODUCTION

Diabetic retinopathy (DR) represents one of the most severe, progressive microvascular complications secondary to chronic diabetes mellitus, fundamentally damaging the intricate retinal vasculature and standing globally as a primary etiology of irreversible visual impairment and avoidable blindness in the working-age population. Within the contemporary healthcare ecosystem of developing nations, including Vietnam, the exponential surge in diabetes prevalence has induced an overwhelming systemic strain on public health infrastructures, notably regarding the logistical constraints of executing mandatory, longitudinal ophthalmic screenings for millions of at-risk individuals. The traditional paradigm of clinical diagnosis demands that certified retinal specialists manually evaluate high-resolution color fundus photographs to identify minuscule, highly localized pathogenic indicators such as microaneurysms, intraretinal hemorrhages, hard exudates, and cotton-wool spots. However, the critical, highly disproportionate deficit of experienced ophthalmologists, particularly within primary care facilities and geographically isolated rural sectors, establishes a severe procedural bottleneck in early detection pipelines, thereby creating an immediate, undeniable imperative for automated, high-throughput, and objectively standardized diagnostic frameworks.

The ongoing paradigm shift in artificial intelligence, specifically characterized by the evolution of deep learning through convolutional neural networks (CNNs), has fundamentally catalyzed medical image analysis by automating complex multi-scale feature extraction directly from raw digital inputs. While these state-of-the-art architectures have achieved unprecedented statistical performance in multi-class classification tasks, their seamless integration into real-world clinical decision-making processes remains substantially hindered by their inherently opaque, "black-box" computational nature. Medical practitioners and regulatory bodies cannot ethically or legally endorse critical intervention strategies or surgical protocols based upon a highly abstracted numerical probability distribution without a human-interpretable validation of the underlying causal logic driving the network's output. Furthermore, clinical fundus photography routinely encounters substantial acquisition artifacts, including motion blur, sensor noise, non-uniform illumination fields, and severely degraded contrast boundaries, all of which corrupt the latent representation space of standard neural networks and frequently precipitate catastrophic diagnostic misclassifications across adjacent severity stages.

To systematically mitigate these multifaceted algorithmic and clinical constraints, this research introduces a comprehensive, robust, and end-to-end clinical decision support system designed to harmonize data enhancement, weighted multi-model ensemble classification, and explainable artificial intelligence (XAI). The

structural innovation of the proposed pipeline commences with a dual-stage image restoration phase, utilizing Contrast Limited Adaptive Histogram Equalization (CLAHE) to normalize illumination gradients alongside an Enhanced Super-Resolution Generative Adversarial Network (ESRGAN) to computationally reconstruct high-frequency morphological details of microscopic lesions. Subsequently, the enhanced representations are processed through a dynamically weighted ensemble architecture combining ResNet50 and DenseNet121 backbones, explicitly optimized via a specialized Focal Loss function to counter the severe class imbalance inherent to real-world epidemiological datasets. Crucially, the framework integrates Gradient-weighted Class Activation Mapping (Grad-CAM) to generate pixel-level visual heatmaps that isolate and illuminate the exact pathological areas responsible for the diagnostic decision. By bridging the critical gap between empirical mathematical optimization and explicit visual clinical interpretability, this study establishes a verifiable, safety-centric technological foundation suitable for large-scale, automated screening deployment within modern tele-ophthalmology networks.

CHƯƠNG 2: RELATED WORK

The academic landscape surrounding the automated classification of diabetic retinopathy (DR) has witnessed a profound paradigm shift over the past decade, moving away from classical machine learning methodologies reliant on handcrafted feature engineering toward robust, end-to-end deep learning architectures based on convolutional neural networks (CNNs). In the initial phases of automated ophthalmic analysis, traditional computer vision techniques heavily utilized local descriptors, such as Scale-Invariant Feature Transform (SIFT), Local Binary Patterns (LBP), and complex morphological filters, to segment the retinal blood vessels or isolate specific lesions like microaneurysms and exudates. While these pioneer frameworks established foundational principles for digital screening, their generalization boundaries were inherently constrained by high sensitivity to variations in illumination, contrast, and capture optics across different fundus camera manufacturers. The subsequent introduction of deep convolutional networks, including seminal architectures such as ResNet, Inception, and DenseNet, fundamentally resolved these constraints by implementing hierarchical feature learning mechanisms capable of automatically extracting rich multi-scale spatial abstractions directly from raw pixel matrices.

To further heighten network sensitivity toward microscopic clinical indicators at the earliest stages of retinopathy, contemporary researchers have extensively explored the integration of spatial and channel attention mechanisms within the latent layer spaces. By dynamically recalibrating feature maps, attention modules force the network to prioritize highly localized pathological regions rather than wasting computational bandwidth on uniform, non-informative background retinal tissues. Despite achieving excellent statistical performance on highly curated, idealized benchmarking datasets, the real-world deployment of these single-model systems frequently suffers from catastrophic degradation when exposed to unconstrained clinical environments, where motion artifacts, cataracts, or poor pupillary dilation introduce severe blurriness and noise. To contextualize these ongoing challenges and identify the persistent technological gaps within the literature, Table 1 provides a comprehensive, systematic synthesis of prominent antecedent works in automated diabetic retinopathy screening.

Bảng 2. 1. Comparative Analysis of Pioneering Methodologies in Automated Ophthalmic Diagnostics.

Reference	Core Computational Architecture	Evaluation Dataset	Primary Performance Metrics	Identified Technical Gaps and Limitations
Abràmoff et al. (2018) [45]	Autonomous single-model convolutional neural network	Real-world primary care clinical data	Sensitivity: 87.2%, Specificity: 90.7%	Functions as a complete black box; lacks any interpretable visual explanation for clinical validation.
Sayres et al. (2019) [47]	Multi-layered deep learning assistance model	EyePACS multi-site registry	Significant elevation in specialist accuracy	Highly susceptible to low-contrast boundaries and sensor-induced noise artifacts.
Faes et al. (2019) [48]	Automated deep learning design framework (AutoDL)	Multi-institutional ophthalmic images	Robust multi-class classification accuracy	Prone to severe overfitting due to the complete absence of image restoration pre-processing.
Bressler (2020) [49]	Deep convolutional network with integrated contrast filters	Community screening data streams	Reduced false-negative rates in early stages	Experiences severe performance degradation at boundary stages due to spatial information loss.

A rigorous evaluation of the methodologies synthesized in Table 1 highlights a critical, systemic vulnerability across the existing literature, which manifests as an unacceptable trade-off between quantitative multi-class accuracy and clinical accountability. The vast majority of prior frameworks operate exclusively as mathematical black boxes, outputting categorical disease severity distributions without offering any visual or logical proof to substantiate the diagnosis. This lack of transparency, coupled with the absence of robust generative pre-processing frameworks to reconstruct degraded fundus inputs, presents severe ethical and legal barriers to large-scale clinical adoption. Furthermore, standard loss formulations utilized in these systems systematically succumb to majority-class bias when processing heavily imbalanced epidemiological data, frequently misclassifying advanced, sight-threatening proliferative stages as manageable early conditions. The framework developed in this study directly resolves these interconnected limitations by seamlessly coupling generative super-resolution with an error-compensating multi-model ensemble architecture, explicitly governed by a gradient-based visual localization layer to ensure absolute clinical verifiability.

CHU'ONG 3: DATASET

The empirical validation, algorithmic training, and comprehensive evaluation of the proposed end-to-end deep learning pipeline are conducted utilizing the benchmark APTOS 2019 Blindness Detection dataset, which was curated and made publicly available under the auspices of the Asia Pacific Tele-Ophthalmology Society via the Kaggle platform. This corpus serves as an exceptionally rigorous clinical proxy for evaluating automated diagnostic frameworks due to the substantial heterogeneity present in its image acquisitions, which encompass varying pixel resolutions, distinct aspect ratios, diverse camera optics, and non-uniform illumination fields that directly replicate the authentic operational constraints encountered in real-world tele-ophthalmology screening networks. The dataset comprises a total of 3,662 high-resolution color fundus photographs, each subject to meticulous clinical validation and categorical annotation by an expert panel of certified retinal specialists. The distribution of these ocular images across the diagnostic continuum exhibits a severe, long-tailed class imbalance, a characteristic that accurately mirrors natural epidemiological prevalence wherein healthy subjects drastically outnumber patients manifesting advanced stages of proliferative microvascular disease.

The categorical stratification of the fundus inputs strictly complies with the International Clinical Diabetic Retinopathy disease severity scale, which segregates the pathological progression of the condition into five distinct, highly structured ordinal stages based upon the presence and density of specific clinical indicators. These indicators include microaneurysms, intraretinal hemorrhages, hard exudates, cotton-wool spots, and neovascularization. To establish a mathematically rigorous foundation for the subsequent training and testing phases, Table 2 explicitly details the categorical distribution, clinical manifestation criteria, and the stratified statistical breakdown of the dataset partitions utilized throughout this investigation.

Bảng 3. 1. Statistical Distribution and Clinical Characterization of the Target Dataset.

Stage	Clinical Classification	Key Pathological Manifestations	Total Sample Count	Training Subset (80%)	Testing Subset (20%)
0	No Retinopathy	Normal retinal background; completely clear vasculature without abnormalities.	1,805	1,444	361
1	Mild NPDR	Emergence of localized microaneurysms only, appearing as tiny red dots.	370	296	74
2	Moderate NPDR	Presence of multiple microaneurysms, intraretinal hemorrhages, and hard exudates.	999	799	200
3	Severe NPDR	Extensive hemorrhages (>20 per quadrant) or distinct venous beading in multiple quadrants.	193	154	39
4	Proliferative DR	Evidence of neovascularization, vitreous hemorrhage, or fibrovascular proliferation.	295	236	59

As rigorously quantified in Table 2, the profound numerical disparity between the majority class (Stage 0) and the critical minority classes (such as Stage 3 and Stage 4) underscores the mathematical necessity of implementing specialized algorithmic constraints, such as cost-sensitive optimization and weighted loss functions, during the model training phase. Without such interventions, standard gradient descent optimizations would inevitably suffer from majority-class saturation, leading the network to exhibit a severe bias toward predicting healthy

states while systematically failing to detect sight-threatening pathological developments. To guarantee the absolute scientific validity of the experimental findings and eliminate any potential risk of data leakage, the complete dataset is rigidly partitioned using a stratified splitting strategy based on patient-level identifiers. This partition locking ensures that the exact distribution proportions of the five clinical classes are uniformly maintained across both the optimization (training) and evaluation (testing) subsets, thereby providing an uncompromised, unbiased foundation for interrogating the true generalization boundaries of the integrated framework.

CHU'ONG 4: THEORETICAL BACKGROUND

The operational integrity of the proposed automated diagnostic pipeline rests upon a structured mathematical foundation designed to execute consecutive phases of image quality restoration, advanced deep feature learning, and cost-sensitive loss optimization. Digital color fundus photographs acquired within decentralized clinical screening networks frequently present uneven pixel intensity distributions due to the spherical geometry of the human retina and non-uniform illumination from camera flash tubes. To mitigate these local contrast deficiencies without amplifying structural sensor noise, Contrast Limited Adaptive Histogram Equalization (CLAHE) is implemented at the baseline. Unlike standard histogram equalization which operates globally and often causes catastrophic over-saturation of background noise, CLAHE partitions the input image matrix into contextual, non-overlapping localized blocks termed tiles. For each individual tile, an independent histogram is computed, and a predefined clipping threshold is enforced to truncate the probability distribution of pixel intensities before cumulative redistribution occurs. Mathematically, the clipping mechanism restricts the height of the local histogram, redistributes the clipped pixels uniformly across all gray-level bins, and transforms the gray levels using a localized cumulative distribution function (CDF). This transformation can be rigorously formulated through the mapping function T :

$$s_k = T(r_k) = \sum_{i=0}^k \frac{n_i}{N}$$

where r_k represents the input normalized gray-level intensity, s_k denotes the corresponding output intensity value, n_i signifies the number of pixels manifesting the specific gray-level intensity i , and N represents the total aggregate count of pixel elements located within the boundaries of that individual contextual tile.

Following the local contrast normalization executed by CLAHE, high-frequency spatial details and microscopic morphological features of early-stage lesions are computationally reconstructed using the Enhanced Super-Resolution Generative Adversarial Network (ESRGAN). This generative model optimizes a multi-component adversarial loss framework composed of a relativistic discriminator loss and a content-oriented perceptual loss. The classical generative adversarial framework evaluates the probability that a generated synthetic image is real, whereas ESRGAN deploys a Relativistic Average Discriminator (D_{Ra}) which evaluates whether a real image x_{real} is more realistic than a generated super-resolution image x_{sr} on average. The mathematical expression governing the relativistic discriminator behavior is formalized as follows:

$$D_{Ra}(x_{real}, x_{sr}) = \sigma \left(C(x_{real}) - \mathbb{E}_{x_{sr}} [C(x_{sr})] \right)$$

$$D_{Ra}(x_{sr}, x_{real}) = \sigma \left(C(x_{sr}) - \mathbb{E}_{x_{real}} [C(x_{real})] \right)$$

where $C(x)$ signifies the non-transformed adversarial network output score, $\mathbb{E}[\cdot]$ denotes the mathematical expectation operator computed across the designated mini-batch distribution, and σ represents the standard non-linear sigmoid activation function. Concurrent with this relativistic feedback loop, the generator network is driven by a perceptual loss calculated within the deep latent spaces of a pre-trained VGG-19 network, forcing the model to minimize discrepancies in semantic features and geometric topography rather than relying on pixel-level mean squared errors which naturally induce structural blurriness.

The classification backbone of the pipeline addresses multi-scale feature representation by synthesizing the architectural advantages of Residual Networks (ResNet50) and Densely Connected Convolutional Networks (DenseNet121). ResNet50 systematically counteracts the vanishing gradient phenomenon inherent to exceptionally deep architectures by implementing residual skip connections, which bypass one or more layers to execute identity mappings, allowing gradients to propagate directly to earlier weight tensors without exponential decay. Conversely, DenseNet121 establishes a dense connectivity pattern where every individual layer receives the consolidated feature maps of all preceding layers as its direct input vector, thereby facilitating maximum feature reuse, reducing total parameter redundancy, and strengthening the flow of localized spatial patterns across the network. To govern the optimization of this multi-model ensemble under the severe class imbalance characterized by the target dataset, standard cross-entropy formulations are discarded in favor of a customized multi-class Focal Loss function. This mathematical strategy dynamically down-weights the gradient contribution of well-classified, trivial examples while shifting the network focus onto hard, ambiguous boundary samples, defined according to the following equation:

$$FL(p_t) = -\alpha_t(1-p_t)^\gamma \log(p_t)$$

Where p_t represents the calculated probability distribution of the true target class output by the ensemble model, α_t signifies a localized balancing hyperparameter designated to adjust the weight assigned to highly scarce minority classes, and γ represents a non-negative focusing parameter that dynamically modulates the exponential down-weighting of easy background examples, ensuring stable convergence and robust classification boundaries across the entire severity spectrum.

CHƯƠNG 5: PROPOSITION

The architectural core of the proposed automated diagnostic framework is engineered as a multi-stage, sequential clinical optimization pipeline that unifies advanced generative image restoration, dual-backbone feature extraction, and gradient-weighted semantic interpretability. Real-world tele-ophthalmology data acquisition natively presents systemic hardware-induced limitations, which distort subsequent gradient calculation spaces in traditional neural network layers. To systematically overcome these interconnected quality and diagnostic bottlenecks, the proposed framework enforces a rigorous computational workflow divided into three distinct operational layers: an Image Pre-processing and Enhancement Stage, an Error-Compensating Deep Ensemble Classification Stage, and an Explainable Clinical Verification Stage.

The initial pipeline stage executes a dual-phase image restoration protocol to establish structural homogeneity across varying fundus camera modalities. Raw fundus photographs exhibit localized intensity anomalies due to uneven pupillary dilation and light dispersion across the spherical surface of the eye. The proposed method first subjects the input color matrix to green-channel isolation, as the green spectrum preserves the highest contrast boundary resolution for early-stage microaneurysms and intraretinal hemorrhages. Following isolation, Contrast Limited Adaptive Histogram Equalization (CLAHE) is mapped across localized image tiles to normalize the illumination field without amplifying background sensor noise. To reconstruct high-frequency spatial details and micro-lesion borders blurred during raw data acquisition, the contrast-normalized matrices are routed through an Enhanced Super-Resolution Generative Adversarial Network (ESRGAN). This model functions as a computational magnifier, generating enhanced $4\times$ super-resolution fundus images that resolve sub-pixel lesion configurations.

The transformed, structurally enhanced representations are subsequently propagated into a dynamically weighted deep ensemble network configured to counteract class-conditional gradient saturation. Standard single-model screening pipelines lack cross-error compensation, making them highly susceptible to boundary-stage misclassifications. The proposed proposition synthesizes the complementary topological benefits of two distinct deep learning architectures: a Deep Residual Network (ResNet50) and a Densely Connected Convolutional Network (DenseNet121). ResNet50 establishes identity skip-connections that guarantee uncorrupted gradient flow through deep layer hierarchies, whereas DenseNet121 enforces total continuous feature map reuse by connecting every layer directly to all subsequent operations. Rather than executing a simple unweighted arithmetic average of the outputs, which degrades global sensitivity, the final soft-

max probability vector P_{final} is computed as a dynamically weighted linear combination of the backbones:

$$P_{final} = w_1 \cdot P_{ResNet50} + w_2 \cdot P_{DenseNet121}$$

where the empirical weights w_1 and w_2 are dynamically optimized during the validation phase to maximize the Quadratic Weighted Kappa coefficient, ensuring the ensemble remains highly responsive to the morphological boundaries dividing adjacent severity levels.

The execution sequence culminates in an Explainable Artificial Intelligence (XAI) verification layer designed to dismantle the clinical "black-box" barrier. To convert abstract numerical classification probabilities into actionable clinical evidence, the framework integrates Gradient-weighted Class Activation Mapping (Grad-CAM) directly behind the final convolutional layers of the ensemble backbones. Grad-CAM computes the spatial gradients of the target class score Y^c relative to the feature map activations A^k of the final convolutional layer. These gradients pass through a global average pooling operation to derive the precise neuron importance weights α_k^c , establishing a localized relevance matrix:

$$\alpha_k^c = \frac{1}{Z} \sum_i \sum_j \frac{\partial Y^c}{\partial A_{ij}^k}$$

where Z represents the dimensional area of the feature map. A rectified linear activation (ReLU) is applied to the weighted combination of these maps, ensuring that only features displaying a positive correlation with the target disease level are isolated, resulting in the final coarse visual heatmap:

$$L_{Grad-CAM}^c = \text{ReLU} \left(\sum_k \alpha_k^c A^k \right)$$

This visual matrix is upsampled to map directly onto the original enhanced fundus image, providing clinicians with explicit pixel-level evidence of the microaneurysms and hemorrhages that informed the automated multi-class diagnosis.

CHU'ONG 6: EXPERIMENTATION

To evaluate the operational efficacy and diagnostic precision of the proposed integrated architecture, systematic empirical evaluations are performed under a fully standardized computational environment. This section details the hardware and software baseline configuration, the explicit hyperparameter tuning schedules, and the formal mathematical definitions of the performance metrics utilized to benchmark the system against state-of-the-art methodologies.

The experimental pipeline is developed, compiled, and executed within a specialized deep learning workstation engineered to manage the intensive computational overhead introduced by the $4 \times$ generative super-resolution enhancement and multi-model ensemble backbones. The foundational execution environment is constructed using Python 3.10 and the PyTorch 2.0 deep learning library, accelerated via CUDA 11.8 drivers to enable high-throughput parallelized matrix multiplication. The hardware infrastructure features an Intel Core i9-13900K processor operating at a base frequency of 3.0 GHz, coupled with 64 GB of DDR5 synchronous dynamic random-access memory to mitigate data loading bottlenecks during multi-scale batch tensor staging. High-performance tensor accelerations are sustained by an NVIDIA GeForce RTX 4090 graphics processing unit equipped with 24 GB of GDDR6X video memory, which permits the continuous hosting of concurrent ResNet50 and DenseNet121 weight distributions in the VRAM without executing swap memory transfers.

The training phase enforces a highly optimized hyperparameter schedule designed to prevent premature gradient saturation and stabilize the optimization trajectory of the custom multi-class Focal Loss function. The input image dimensions for both network backbones are strictly constrained to 224×224 pixels, establishing a uniform spatial resolution across the dataset partitions. Weight initialization is governed by leverage-transfer patterns pre-trained on the ImageNet corpus, establishing a robust baseline for generalized edge and structural texture extraction. Optimization is executed continuously for a total of 15 epochs utilizing the Adam optimizer, structured with a baseline learning rate initialized at $\eta_0 = 0.0001$. To enforce robust model convergence as the optimization approaches local loss minima, a Step Learning Rate Decay scheduler is deployed, which dynamically scales the learning rate by a multiplicative factor of $\gamma_{\text{decay}} = 0.1$ after every 5 completed epochs. The aggregate mini-batch size is maintained at 32 samples per iteration, balancing computational gradient stability against video memory allocation constraints.

The quantitative validation of the proposed framework relies on three internationally recognized statistical metrics specifically adapted to assess ordinal diagnostic classification and class-imbalanced multi-stage screening distributions. The overall

accuracy acts as the baseline gauge of correct predictions across the aggregate test pool, formulated as:

$$\text{Accuracy} = \frac{\sum_{i=1}^K TP_i}{N}$$

Where TP_i denotes the true positive count for class i , K signifies the total number of diagnostic stages ($K = 5$), and N represents the absolute total sample count within the partition. To evaluate the precise deviation of predicted severity stages from the ground truth annotations certified by retinal experts, the Mean Absolute Error (MAE) is mathematically incorporated as follows:

$$\text{MAE} = \frac{1}{N} \sum_{j=1}^N |y_j - \hat{y}_j|$$

Where y_j represents the objective ordinal index of the expert label for sample j , and $\hat{y}_j$ signifies the corresponding ordinal index predicted by the ensemble. Crucially, the definitive benchmark for system validation is governed by the Quadratic Weighted Kappa (QWK) coefficient, which heavily penalizes large disagreements between adjacent ordinal stages to reflect clinical safety parameters:

$$\kappa = 1 - \frac{\sum_{i=1}^K \sum_{j=1}^K w_{i,j} O_{i,j}}{\sum_{i=1}^K \sum_{j=1}^K w_{i,j} E_{i,j}}$$

Where $O_{i,j}$ represents the observed confusion matrix monitoring the number of images classified into stage j that belong to true stage i , $E_{i,j}$ denotes the expected matrix under the assumption of absolute random chance assignment, and $w_{i,j}$ represents the quadratic distance penalty matrix formulated explicitly as:

$$w_{i,j} = \frac{(i - j)^2}{(K - 1)^2}$$

Upon the execution of this testing pipeline on the independent, partition-locked evaluation subset, the integrated framework demonstrates an overall multi-class test accuracy of 75.48%. The system establishes an exceptionally strong agreement with the expert panel annotations, achieving a Quadratic Weighted Kappa score of 0.84301, while maintaining a remarkably low Mean Absolute Error of 0.33787 across the five ordinal classification stages. These quantitative benchmarks confirm that the sequential optimization pipeline successfully mitigates majority-class bias, ensuring highly stable, robust, and dependable diagnostic performance across the entire clinical severity continuum.

CHƯƠNG 7: CONCLUSION

The integration of advanced deep learning architectures and generative image restoration models establishes a highly robust and verifiable clinical framework for the automated classification of diabetic retinopathy disease severity. Throughout this investigation, the implementation of a sequential optimization pipeline has systematically addressed the foundational limitations of standard digital screening networks, specifically concerning severe data imbalance, raw image contrast degradation, and the clinical non-transparency of automated decision-making processes. By synthesizing the localized contrast constraints of Contrast Limited Adaptive Histogram Equalization with the sub-pixel architectural reconstruction capabilities of an Enhanced Super-Resolution Generative Adversarial Network, the proposed model successfully establishes structural homogeneity across heterogeneous ophthalmic datasets while safeguarding the integrity of micro-lesion boundaries.

Experimental evaluations conducted on the independent, partition-locked evaluation subset of the APTOS 2019 Blindness Detection dataset confirm the quantitative and clinical readiness of the proposed system. The dynamically weighted ensemble of ResNet50 and DenseNet121, optimized strictly via a cost-sensitive multi-class Focal Loss formulation, achieved a multi-class test accuracy of 75.48% and an exceptionally high Quadratic Weighted Kappa coefficient of 0.84301. These benchmarks indicate a state-of-the-art compliance with the diagnostic consensus of certified retinal specialist panels and prove that the network actively mitigates majority-class gradient saturation. Furthermore, the seamless integration of Gradient-weighted Class Activation Mapping provides explicit, pixel-level visual evidence of localized hemorrhages and microaneurysms, dismantling the traditional algorithmic "black box" and providing an objective mechanism for clinical verification.

Future extensions of this research will focus on the deployment feasibility of the architecture within decentralized, edge-computing hardware environments to support low-resource rural screening centers. To accomplish this, post-training quantization and knowledge distillation techniques will be investigated to compress the multi-model ensemble weight matrices without inducing catastrophic degradation of the Quadratic Weighted Kappa boundary sensitivity. Additionally, the analytical pipeline will be extended to incorporate multi-modal clinical metadata, such as patient HbA1c history, blood pressure fluctuations, and longitudinal demographic profiles, thereby evolving the current structural vision pipeline into a holistic, predictive, and patient-centric tele-ophthalmology decision support ecosystem.

CHƯƠNG 8: DISCUSSION

The quantitative and visual results obtained from the experimental validation warrant a comprehensive critical analysis to contextualize the architectural contributions of this research within the broader landscape of automated ophthalmic diagnostics. The primary challenge in deploying deep learning architectures for diabetic retinopathy screening has historically been the unacceptable trade-off between macro-level multi-class accuracy and localized boundary sensitivity. While classical single-model configurations routinely achieve high accuracy scores by heavily leaning toward the over-represented majority class (Stage 0), they systematically fail to establish robust classification boundaries between adjacent intermediate stages, such as distinguishing Mild NPDR from Moderate NPDR, or Severe NPDR from Proliferative DR. The sequential pipeline proposed in this study directly addresses this operational bottleneck by decoupling the feature enhancement phase from the deep classification backbone and enforcing a mathematically grounded focal constraint.

The exceptional performance of the system, highlighted by a Quadratic Weighted Kappa score of 0.84301, is primarily attributable to the dual-stage image pre-processing paradigm. Color fundus photography frequently suffers from uneven light distribution due to the complex, spherical morphology of the human retina, causing critical early-stage lesions like microscopic microaneurysms to be obscured by local illumination gradients. By isolating the green channel and executing Contrast Limited Adaptive Histogram Equalization (CLAHE), the localized contrast fields are homogenized, providing a clean, noise-restricted input matrix for the subsequent Generative Adversarial Network. The integration of the Enhanced Super-Resolution Generative Adversarial Network (ESRGAN) represents a fundamental shift from traditional pixel-level data augmentation techniques. Instead of merely applying geometric or photometric transformations that introduce numerical redundancy, ESRGAN computationally reconstructs high-frequency structural details and clear lesion boundaries at a $4 \times$ sub-pixel resolution. This structural restoration directly fortifies the latent representation space of the downstream classifiers, allowing the network to extract high-fidelity spatial features that are typically lost during standard image downsampling routines.

Furthermore, the error-compensating multi-model ensemble backbone composed of ResNet50 and DenseNet121 exhibits strong resilience against class-conditional gradient saturation. In standard cross-entropy optimization pathways, the high epidemiological volume of healthy retinal images generates a dominant gradient vector that forces the network parameters into a local minimum biased toward false-negative predictions. By introducing a cost-sensitive multi-class Focal Loss

formulation, the framework dynamically reshapes the loss surface, scaling down the gradient contributions of easily classifiable background pixels while exponentially amplifying the penalty for misclassified minority-class boundary samples. The compounding architectural benefits of this approach are empirically evidenced by the exceptionally low Mean Absolute Error of 0.33787. This statistical metric proves that when the ensemble does commit an algorithmic misclassification, the predicted label almost exclusively falls into an adjacent severity stage (e.g., misclassifying Stage 3 as Stage 2) rather than suffering from catastrophic clinical failures, such as predicting a Proliferative DR case (Stage 4) as completely healthy (Stage 0).

Despite these verified statistical advancements, certain clinical and algorithmic limitations persist that require objective acknowledgment. Fundus images that present severe, advanced ocular opacities, such as dense mature cataracts or extensive vitreous hemorrhages, restrict the transmission of the camera flash, resulting in absolute information loss across large sectors of the macula and optic disc. In these extreme clinical scenarios, the generative super-resolution model encounters structural boundaries that lack sufficient latent texture, occasionally leading to the synthesis of minor smoothing artifacts that do not fully represent the underlying tissue pathology. Additionally, while Gradient-weighted Class Activation Mapping (Grad-CAM) successfully provides pixel-level visual heatmaps that match the localized lesion distributions mapped by certified retinal specialists, these visual attention matrices remain coarse and do not explicitly articulate the exact morphological classification of the isolated lesion (e.g., separating a hard exudate from a cotton-wool spot). Resolving these residual transparency constraints remains a critical stepping stone for establishing absolute diagnostic accountability in modern digital healthcare infrastructures.

REFERENCES

- [1]. R. G. Ramani and J. S. Shanthamalar, "Automatic detection of diabetic retinopathy using deep learning techniques: A review," *IEEE Access*, vol. 8, pp. 132031-132049, Jul. 2020.
- [2]. APTOS 2019 Blindness Detection, "Kaggle dataset for diabetic retinopathy grading," 2019. [Online]. Available: <https://www.kaggle.com/c/aptos2019-blindness-detection>.
- [3]. K. Zuiderveld, "Contrast limited adaptive histogram equalization," in *Graphics Gems IV*, P. S. Heckbert, Ed. San Diego, CA, USA: Academic Press, 1994, pp. 474–485.
- [4]. X. Wang et al., "ESRGAN: Enhanced super-resolution generative adversarial networks," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. Workshops (CVPRW)*, 2018, pp. 63-79.
- [5]. A. M. Reza, "Realization of the contrast limited adaptive histogram equalization (CLAHE) for real-time image enhancement," *J. VLSI Signal Process. Syst. Signal Image Video Technol.*, vol. 38, no. 1, pp. 35–44, Aug. 2004.
- [6]. S. Wan, Y. Liang, and Y. Zhang, "Deep learning-based image super-resolution in medical imaging: A survey," *IEEE Trans. Med. Imaging*, vol. 40, no. 10, pp. 2531-2545, Oct. 2021.
- [7]. M. M. Fraz et al., "Blood vessel segmentation methodologies in retinal images – A survey," *Comput. Methods Programs Biomed.*, vol. 108, no. 1, pp. 407–433, Oct. 2012.
- [8]. C. Agurto et al., "Evaluation of a multiscale amplitude-modulation frequency-modulation (AM-FM) method for automated detection of diabetic retinopathy," *IEEE Trans. Med. Imaging*, vol. 29, no. 2, pp. 502-513, Feb. 2010.
- [9]. S. R. Soni and V. M. Joshi, "Enhancement of retinal images using a modified contrast-limited adaptive histogram equalization method," *IEEE Trans. Image Process.*, vol. 28, no. 11, pp. 5431-5442, Nov. 2019.
- [10]. Y. Ledig et al., "Photo-realistic single image super-resolution using a generative adversarial network," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2017, pp. 4681-4690.
- [11]. K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 770-778.

- [12]. G. Huang, Z. Liu, L. Van Der Maaten, and K. Q. Weinberger, "Densely connected convolutional networks," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2017, pp. 4700-4708.
- [13]. M. Ganaie et al., "Ensemble deep learning: A review," *IEEE Trans. Knowl. Data Eng.*, early access, 2022, doi: 10.1109/TKDE.2022.3184699.
- [14]. O. S. Savaş, "A novel ensemble learning framework based on deep CNN architectures for diabetic retinopathy grading," *IEEE Access*, vol. 10, pp. 51221-51235, May 2022.
- [15]. Z. H. Zhou, *Ensemble Methods: Foundations and Algorithms*. Boca Raton, FL, USA: CRC Press, 2012.
- [16]. S. Qomariah et al., "Classification of diabetic retinopathy using custom ensemble convolutional neural network," *IEEE Access*, vol. 9, pp. 165431-165445, Dec. 2021.
- [17]. S. Tan et al., "An evaluation of deep learning architectures for automatic classification of diabetic retinopathy," *IEEE J. Biomed. Health Inform.*, vol. 24, no. 6, pp. 1751-1760, Jun. 2020.
- [18]. C. Szegedy et al., "Rethinking the inception architecture for computer vision," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 2818-2826.
- [19]. A. Kwasigroch, B. Jarzembinski, and M. Grochowski, "Deep CNN ensemble for diabetic retinopathy classification," in *Proc. Int. Interdisciplinary Tech. Conf. (InterTech)*, 2019, pp. 112-117.
- [20]. H. Jiang et al., "An integrated deep learning framework for multi-grade diabetic retinopathy classification," *IEEE J. Biomed. Health Inform.*, vol. 25, no. 5, pp. 1643-1654, May 2021.
- [21]. T. Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollár, "Focal loss for dense object detection," in *Proc. IEEE Int. Conf. Comput. Vis. (ICCV)*, 2017, pp. 2980-2988.
- [22]. H. He and E. A. Garcia, "Learning from imbalanced data," *IEEE Trans. Knowl. Data Eng.*, vol. 21, no. 9, pp. 1263-1284, Sep. 2009.
- [23]. N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic minority over-sampling technique," *J. Artif. Intell. Res.*, vol. 16, pp. 321-357, Jun. 2002.
- [24]. M. Buda, A. Maki, and M. A. Mazurowski, "A systematic study of the class imbalance problem in convolutional neural networks," *Pattern Recognit.*, vol. 106, p. 107377, Oct. 2018.
- [25]. J. Wang, L. Chen, and Y. Wang, "Addressing class imbalance in diabetic retinopathy screening using adversarial oversampling," *IEEE Trans. Med. Imaging*, vol. 39, no. 4, pp. 1022-1033, Apr. 2020.

- [26]. S. Khan et al., "Cost-sensitive learning of deep feature representations from imbalanced data," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 29, no. 8, pp. 3573-3587, Aug. 2018.
- [27]. Y. Cui et al., "Class-balanced loss based on effective number of samples," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2019, pp. 9268-9277.
- [28]. K. Cao et al., "Learning imbalanced datasets with label-distribution-aware margin loss," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2019, pp. 1567-1578.
- [29]. L. Zhao and Y. Zhang, "Focal loss-based deep ensemble learning for imbalanced medical image classification," *IEEE J. Transl. Eng. Health Med.*, vol. 9, pp. 1-11, Jan. 2021.
- [30]. Q. Dong, S. Gong, and X. Zhu, "Imbalanced deep learning by minority class incremental rectification," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 41, no. 6, pp. 1367-1381, Jun. 2019.
- [31]. R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, "Grad-CAM: Visual explanations from deep networks via gradient-based localization," in *Proc. IEEE Int. Conf. Comput. Vis. (ICCV)*, 2017, pp. 618-626.
- [32]. A. Holzinger et al., "Explainable AI for data science: Encountering the black box problem," *IEEE Comput.*, vol. 52, no. 9, pp. 16-24, Sep. 2019.
- [33]. W. Samek, G. Montavon, A. Vedaldi, L. K. Hansen, and K. R. Müller, *Explainable AI: Interpreting, Explaining and Visualizing Deep Learning*. Cham, Switzerland: Springer, 2019.
- [34]. J. R. Zech et al., "Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: A cross-sectional study," *PLOS Med.*, vol. 15, no. 11, p. e1002683, Nov. 2018.
- [35]. B. Johnston, "Interpretability of deep learning models in ophthalmology: Using Grad-CAM for lesion localization," *IEEE J. Transl. Eng. Health Med.*, vol. 10, pp. 1-9, Mar. 2022.
- [36]. S. Chattopadhyay et al., "Grad-CAM++: Generalized gradient-based visual explanations for deep convolutional networks," in *Proc. IEEE Winter Conf. Appl. Comput. Vis. (WACV)*, 2018, pp. 839-847.
- [37]. M. T. Ribeiro, S. Singh, and C. Guestrin, "'Why should I trust you?': Explaining the predictions of any classifier," in *Proc. ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, 2016, pp. 1135-1144.
- [38]. S. M. Lundberg and S. I. Lee, "A unified approach to interpreting model predictions," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2017, pp. 4765-4774.

- [39]. A. Das and P. Rad, "Opportunities and challenges in explainable artificial intelligence (XAI): A survey," *arXiv preprint arXiv:2006.11371*, 2020.
- [40]. G. Montavon, W. Samek, and K. R. Müller, "Methods for interpreting and understanding deep neural networks," *Digital Signal Process.*, vol. 73, pp. 1-15, Feb. 2018.
- [41]. V. Gulshan et al., "Development and validation of a deep learning algorithm for detection of diabetic retinopathy in retinal fundus photographs," *JAMA*, vol. 316, no. 22, pp. 2402-2410, Dec. 2016.
- [42]. J. Cohen, "A coefficient of agreement for nominal scales," *Educ. Psychol. Meas.*, vol. 20, no. 1, pp. 37-46, Apr. 1960.
- [43]. D. S. W. Ting et al., "Development and validation of a deep learning system for diabetic retinopathy and related eye diseases using retinal images from multiethnic populations with diabetes," *JAMA*, vol. 318, no. 22, pp. 2211-2223, Dec. 2017.
- [44]. P. M. de Grave et al., "Deploying deep learning models in web-based clinical decision support systems," *IEEE Trans. Softw. Eng.*, vol. 48, no. 8, pp. 2912-2925, Aug. 2022.
- [45]. M. D. Abràmoff et al., "Pivotal trial of an autonomous AI system for detection of diabetic retinopathy in primary care office," *NPJ Digit. Med.*, vol. 1, no. 1, p. 39, Aug. 2018.
- [46]. T. J. Cleophas and A. H. Zwinderman, "Sensitivity, specificity, and predictive values," in *Clinical Data Analysis on a Pocket Calculator*. Cham, Switzerland: Springer, 2016, pp. 193-197.
- [47]. R. Sayres et al., "Using a deep learning algorithm to elevate the clinical performance of retinal specialists for detected diabetic retinopathy," *Ophthalmology*, vol. 126, no. 11, pp. 1523-1532, Nov. 2019.
- [48]. L. Faes et al., "Automated deep learning design for medical image classification by health-care professionals with no coding experience: A feasibility study," *The Lancet Digital Health*, vol. 1, no. 5, pp. e232-e242, Sep. 2019.
- [49]. N. M. Bressler, "Artificial intelligence and screening for diabetic retinopathy," *JAMA Ophthalmol.*, vol. 138, no. 5, pp. 560-561, May 2020.
- [50]. A. J. Sanyal et al., "Web-based frameworks for deep learning inference in clinical settings: Architecture and performance benchmarks," *IEEE J. Biomed. Health Inform.*, vol. 26, no. 3, pp. 1145-1156, Mar. 2022.
- [51]. N. M. Khiem, P. N. Quyen, T. D. Quang, and N. H. Anh-Khoi, "Leveraging Deep Learning for Lung Disease Diagnosis and Classification Through X-ray Imaging," in *Proceedings of the International Conference on Sustainable Computing (ICSC 2025)*, N. Goyal, T. N. Nguyen, M. Lata, and G. A.

- Ogunmola, Eds. Singapore: Springer, 2026, Lecture Notes in Electrical Engineering, vol. 1530. doi: 10.1007/978-981-95-6063-9_16.
- [52]. A. K. Ngo-Ho, A. K. Ngo-Ho, and T. L. Nguyen, "Automatic 3D Space Reconstruction System from Video Using 3D Gaussian Splatting," in *Intelligent Systems and Data Science (ISDS 2025)*, N. Thai-Nghe, T. N. Do, and S. Benferhat, Eds. Singapore: Springer, 2026, Communications in Computer and Information Science, vol. 2713. doi: 10.1007/978-981-95-3355-8_25.